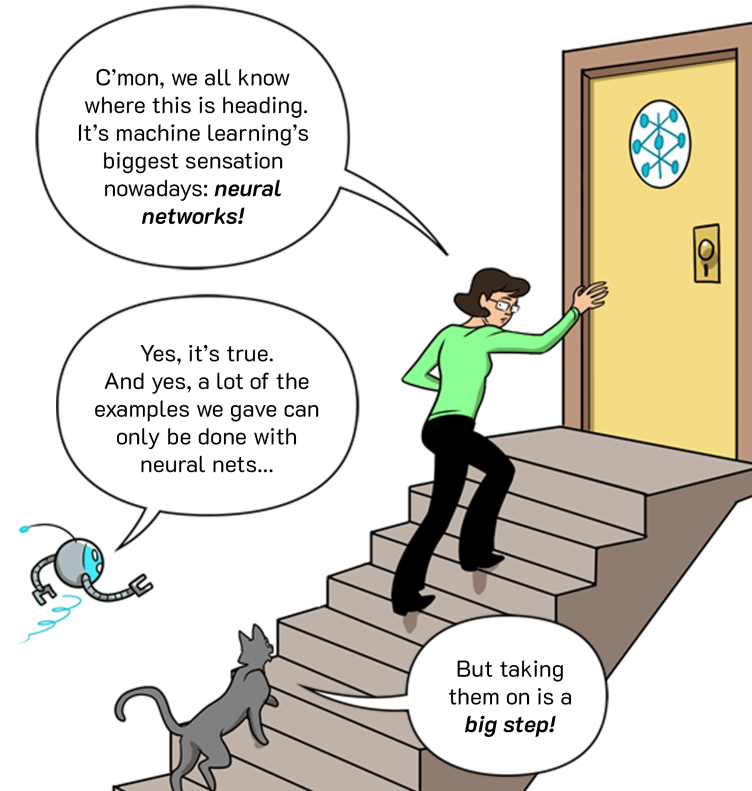


# Introduction to Artificial Neural Networks Part II

Dr. Andrea Santamaria Garcia, Chenran Xu  
Institute of Beam Physics and Technology (KIT)



# Recap from previous lecture

## Neural networks:

- are **powerful function approximators** that are **computationally efficient** for big data.
- rely on **large quantities of training data**, and their performance is affected by the **quality** and **variety** of the data.
- are part of **narrow AI**, which means they are specialized to solve particular tasks and do not generalize to different tasks.
- are used in **supervised** and **unsupervised** learning.
- are parametrized by their **weights and biases**, which are iteratively optimized by minimizing the error between the predicted value and target value, called **loss function**.
- are widely used with **first-order gradient methods** to minimize the loss function.
- update their weights through **backpropagation**, which is based on the multivariate chain rule.

# Important limitation of gradient descent

How much computational time does it take to calculate the gradients?

- Number of points =  $n$
- Number of parameters =  $p$

$$\sum_{i=1}^n \underbrace{(h_w(x_i) - y_i)^2}_{n \text{ terms}} = J(W)$$

$$W \leftarrow W - \alpha \underbrace{\nabla J(W)}_{\substack{\text{Calculate} \\ \text{gradient } p \text{ times}}}$$

Let's consider a small example:

- 2000 data points
- 5-64-64-2 network (578 parameters)
  - ~1.1 million computations for one model update

Evaluates the loss over the entire dataset



calculate the gradient using just a random small part of the observations instead of all of them



**Stochastic gradient descent**

# Stochastic gradient descent

In stochastic gradient descent (SGD) the gradient is approximated by a gradient at a single sample:

repeat until approx.  
minimum

Randomly shuffle samples in the data set  
for  $i = 1, \dots, n$  do:

$$w \leftarrow w - \alpha \nabla J(w_k)$$

## Mini-batch SGD a compromise between GD and SGD


The dataset can be sliced in random mini-batches that are mutually independent.

Several passes can be made over the training set until the algorithm converges.


- A new hyperparameter to tune appears: the minibatch size.

Why is this a good idea, computational efficiency aside?

- In SGD the loss is approximated over a subset of the data, which greatly improves computational efficiency.
- This approximation results in a noisy loss function, different from the “all data” loss function.
- The stochasticity might help in some cases to avoid local minima.

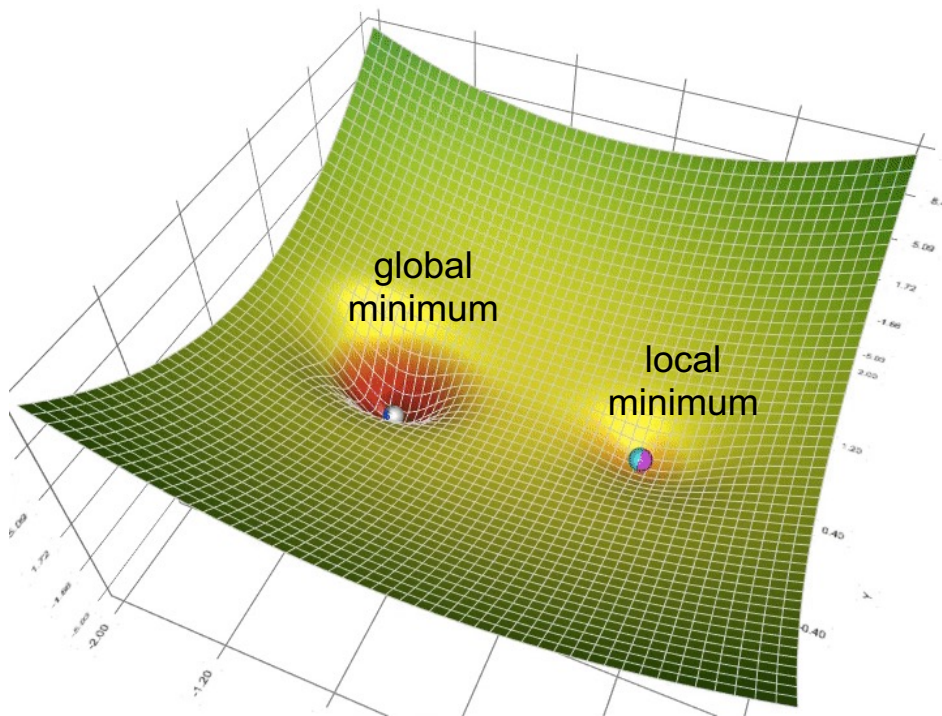


GD



SGD

# Other optimizers vs gradient descent



Animation from [Lili Jiang](#)

## Gradient based methods (first order):

- Gradient descent
- momentum
- **AdaGrad**
- RMSProp
- Adam

} SGD with adaptive learning rates

## Second order gradient methods:

- They provide information about the curvature of the loss function (e.g. Newton's method)
- Complex and difficult to implement
- Expensive in iteration cost and memory occupation
- Active area of research

## Gradient-free methods:

- Genetic algorithms, particle swarm optimization.
- "Neuroevolution"

# Activation functions

The choice of activation function has a large impact on the capability and performance of the neural network!

*They manage the flow of data through the network by activating or deactivating neurons based on their output*

Typically:

- All hidden layers use the same activation function.
- The output layer will use a different activation function since is dependent upon the type of prediction required by the model.
- Activation functions are differentiable (first-order derivative can be calculated for a given input value).

## Hidden layers

- ReLU (all NNs)
- Tanh (RNNs)
- Sigmoid (RNNs)


## Output layers

- Linear (regression)
- Sigmoid (binary classification)
- Softmax (multiclass classification)

# Activation functions


The choice of activation function has a large impact on the capability and performance of the neural network!

## Sigmoid / Logistic




- Maps the input to a  $[0, 1]$  range, useful to predict probabilities.
- The smaller the input, the closer it will be to 0.
- Saturates at the tail of 0 and 1.
- Output is not zero centered and will always be the same sign and in a small range, which makes training more difficult and unstable.

## Tanh



- Maps the input to a  $[-1, 1]$ .
- The smaller the input, the closer it will be to -1.
- The output is zero centered, so output values can be easily mapped to strongly positive, negative or neutral.
- Neurons saturate for large negative and positive values.

## ReLU




- Only positive values pass through (neuron deactivated if output is negative).
- Very computationally efficient.
- Non saturating property accelerates GD convergence.
- The gradient is also zero for negative values, which can create dead neurons that never get activated (leaky ReLU).

# Vanishing gradients

appear in backpropagation using gradient-based methods in deep networks

not good in hidden layers




Maximum of gradient 0.25  
With chain rule the gradient product can become very small

$$\frac{\partial J(w)}{\partial w^{(0)}} = \frac{\partial J(w)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a^{(1)}} \dots \frac{\partial a^{(l)}}{\partial w^{(0)}}$$

$$0.2 \times 0.15 \times 0.22 \times 0.09 \dots$$

No response to changes in input


Very narrow range, small values



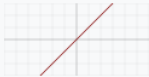


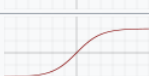

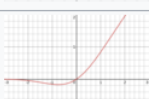




When the partial derivative vanishes the weights are not updated anymore

$$w \leftarrow w - \alpha \nabla J(w_k)$$

good for hidden layers



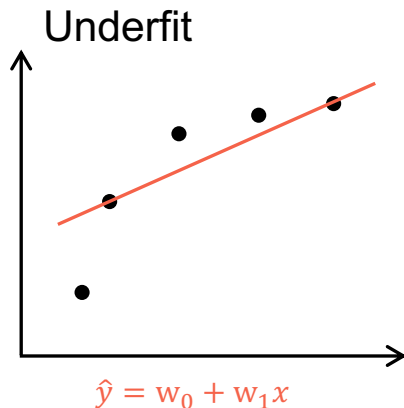




Name	Plot	Function, $g(x)$	Derivative of $g$ , $g'(x)$
Identity		$x$	1
Binary step		$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	$\begin{cases} 0 & \text{if } x \neq 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$
Logistic, sigmoid, or soft step		$\sigma(x) \doteq \frac{1}{1 + e^{-x}}$	$g(x)(1 - g(x))$
Hyperbolic tangent (tanh)		$\tanh(x) \doteq \frac{e^x - e^{-x}}{e^x + e^{-x}}$	$1 - g(x)^2$
Rectified linear unit (ReLU) <sup>[8]</sup>		$(x)^+ \doteq \begin{cases} 0 & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$ $= \max(0, x) = x \mathbf{1}_{x>0}$	$\begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$
Gaussian Error Linear Unit (GELU) <sup>[5]</sup>		$\frac{1}{2}x \left( 1 + \operatorname{erf} \left( \frac{x}{\sqrt{2}} \right) \right)$ $= x\Phi(x)$	$\Phi(x) + x\phi(x)$
Softplus <sup>[9]</sup>		$\ln(1 + e^x)$	$\frac{1}{1 + e^{-x}}$
Exponential linear unit (ELU) <sup>[10]</sup>		$\begin{cases} \alpha(e^x - 1) & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$ with parameter $\alpha$	$\begin{cases} \alpha e^x & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ 1 & \text{if } x = 0 \text{ and } \alpha = 1 \end{cases}$
Scaled exponential linear unit (SELU) <sup>[11]</sup>		$\lambda \begin{cases} \alpha(e^x - 1) & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$ with parameters $\lambda = 1.0507$ and $\alpha = 1.67326$	$\lambda \begin{cases} \alpha e^x & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$
Leaky rectified linear unit (Leaky ReLU) <sup>[12]</sup>		$\begin{cases} 0.01x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	$\begin{cases} 0.01 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$

# Overfitting

- Overfitting happens when the NN learns too many details about the training data and will fail to generalize to unseen data.
- This happens due to:
  - Overly complex models with too many parameters (e.g. deep neural networks).
  - Training for too long, capturing noise instead of genuine patterns.



# Regularization

Improves generalization on unseen data by constraining the optimization problem to discourage complex models

## Early stopping:

- Stop training before overfitting.
- Criteria: when the loss does not improve beyond a certain threshold.
- The “patience” parameter allows you to specify how much to wait after the threshold before stopping training.

## Dropout:

- Randomly deactivate neurons from the during training in each iteration (equivalent to training different NNs).
- Reduces number of parameters.
- Avoids relying on certain nodes that only learn certain patterns.

## Weight regularization

- Adds a weight penalty term to the loss function, penalizing large weights (too sensitive, large changes in the output).
- The term increases the error, forcing the network to minimize the weights contributing more to the loss.

## Batch normalization:

- Normalizes the inputs to layers across each mini-batch to reduce internal covariate shift (the distribution of each layer's inputs changes during training, as the parameters of the previous layers change, slowing down training).

# Stages in supervised learning

## Training Phase:

Learns the basic mapping between input and output  
You can develop several models

- The algorithm is trained using a labeled dataset consisting of training examples  $(x, y)$ .
- The model makes predictions on the training data based on the current state of its parameters.
- A **loss function** is used to measure the difference between the model's predictions and the actual target values for the training data.
- The objective of the training is to minimize this loss function. This is done using **optimization algorithms** like **gradient descent**.

50% of data  
Curated  
Representative  
“Gold standard”

# Stages in supervised learning

## Validation Phase:

Select best performing model or approach

- The model's performance is evaluated on a separate dataset not seen by the model during training (validation dataset).
- This phase helps in tuning the model's hyperparameters and provides an estimate of how well the model has generalized to unseen data.

## Testing Phase:

Evaluation of final model performance

- Once the model is trained and validated, its performance is tested on another set of unseen data (test dataset).
- This phase provides an unbiased evaluation of the final model fit on the training dataset.

25% of data

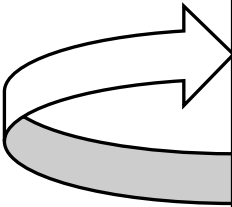
Real-world data  
Data of your study

25% of data

Real-world data  
Data of your study

Not used in training for  
unbiased performance  
estimation

# Summary: how to train a neural network

- 
1. **Select data features + perform data scaling**
  2. Choose network architecture
  3. Choose activation functions
  4. Choose loss function & convergence criteria
  5. Choose an optimizer & set its hyperparameters
  6. Choose number of epochs to train
  7. Decide on regularization techniques
  8. **Perform forward propagation**
  9. **Compute gradients with backwards propagation**
  10. **Update weights & keep track of loss**
  11. **Evaluate the model's efficiency**

# Supervised learning training loop

## Pseudo code for training a NN with SGD with mini-batches

```
randomly initialize the NN weights  $\mathbf{w}$ 
for  $i = 1, \dots, n_{\text{epoch}}$  do:
    randomly shuffle samples in the data set
    for  $batch$  in  $mini\text{-}batches$  do: # loop over the data set in batches
        perform forward pass  $\hat{y} = f(x)$ 
        calculate loss  $J(\hat{y}, y)$ 
        update weights  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla_{\mathbf{w}} J$  # one step SGD update
```

- $n_{\text{epoch}}$  is the number of times the NN goes over the entire data set
- The data set is randomly shuffled and split into small chunks (mini-batches) each with  $batch\_size$  samples
- One batch is fed through the the NN (in parallel) and weights are updated once.

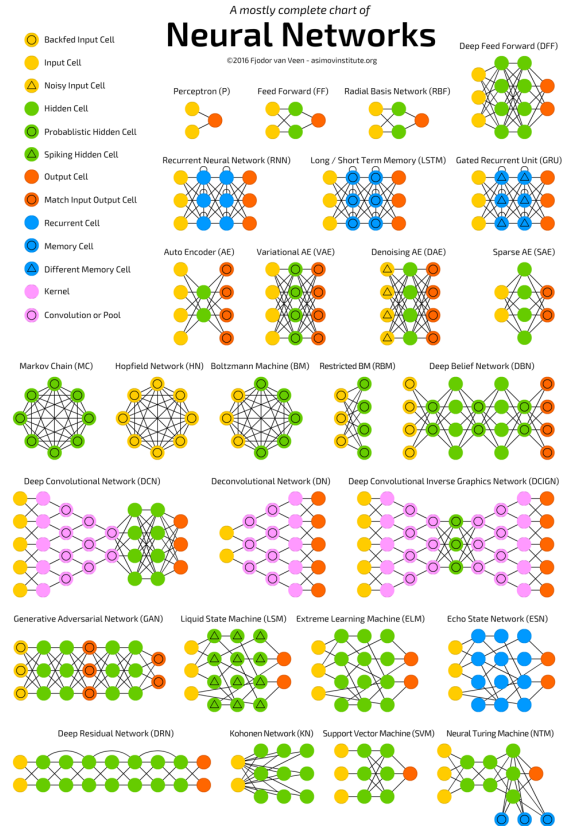
# Further considerations

- **Error analysis:** under which conditions does the model perform poorly?
- **Statistical testing:** are the performance metrics between models statistically significant?
- **Robustness and generalization:** how well does the model perform on variations of the test data it was not explicitly trained on?
- **Real-world performance:** evaluate the model in a real-world setting or a simulation that closely mimics the production environment.
- **Resource utilization:** assess the model's resource usage, like inference time, memory footprint, and power consumption.



**A sneak peak of more advanced  
concepts:  
deep neural networks**

# Neural network architectures




- Until now we only looked at the simple feed-forward, fully-connected case.
- Neural networks can have very different architectures.

# Autoencoder (AE)

- Feed-forward structure, used for different purposes
- **Encodes** the input to low-dimensional latent space, and **decodes** to the original shape

## Use cases

- Feature extraction
- Denoising input data
- Generative models



# Convolutional Neural Network (CNN)

- Used for **image-related** tasks
- Applies spatial convolutions to the input
- Translation invariance
- Hierarchical features: deeper layers detect more complicated features
- Less parameters needed than fully-connected structure

LeNet-5 structure, Y. Lecun (1998). 61k parameters






Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

ResNet-50, K. He (2015). 25M parameters

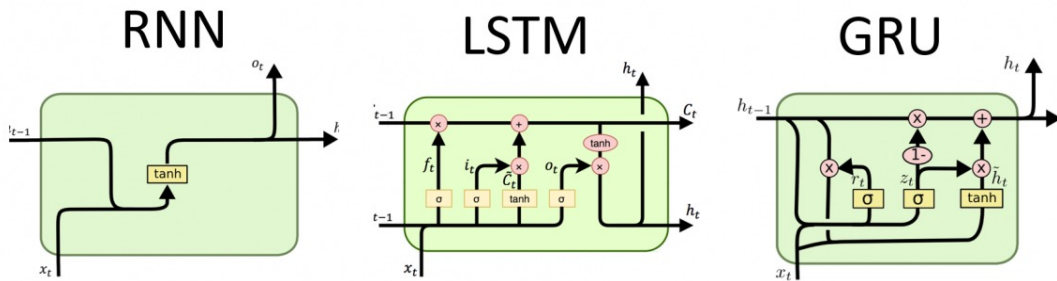


# Recurrent Neural Network (RNN)

- Used in **time-series** data.  
e.g. Natural language processing (NLP), forecasting, ...
- Contains hidden state variables (memory, context).
- Allows variational lengths of inputs and outputs.



<https://www.geeksforgeeks.org/introduction-to-recurrent-neural-network/>




<http://dprogrammer.org/rnn-lstm-gru>

# Transformer

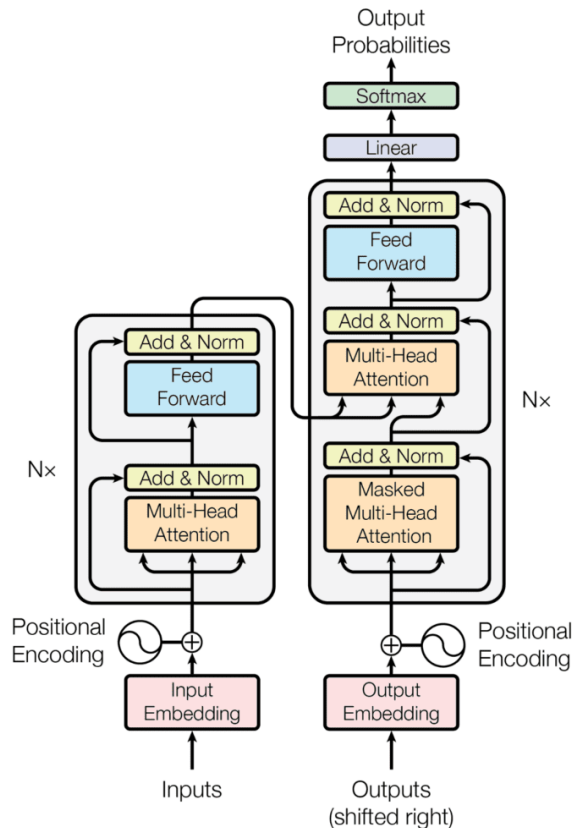

[Attention is all you need, 2017](#)

- Probably the most successful NN structure nowadays
- **Architecture behind the LLMs (ChatGPT,...)**
- Use only the *attention* mechanism without the recurrent structure
- Parallelizable -> faster training

Scaled Dot-Product Attention



Multi-Head Attention



# There is not one library to rule them all



## Neural networks/ Deep learning



Tensorflow backend



## ML algorithms / optimization

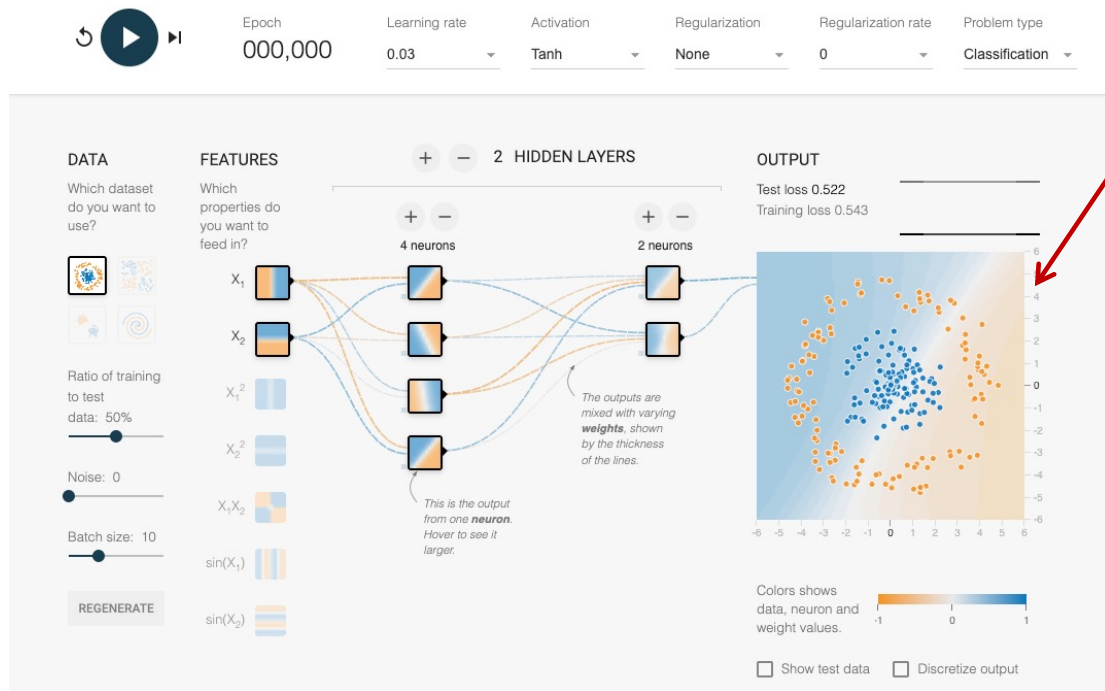


+



# Let's put everything we have learned to practice!

Go to: <https://playground.tensorflow.org/>



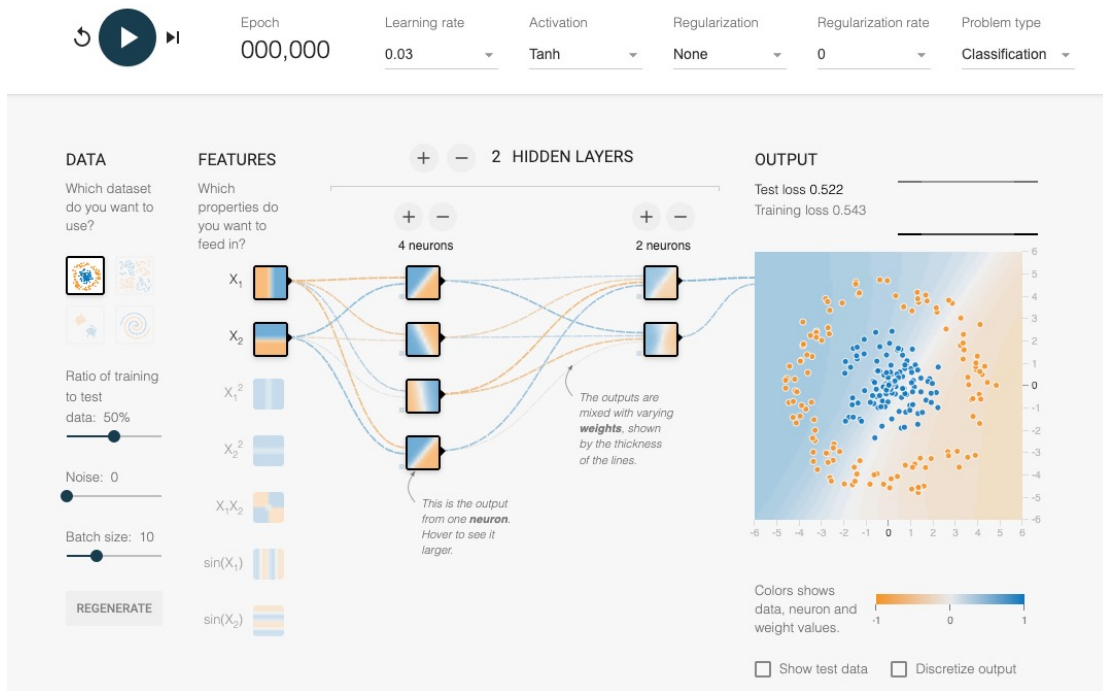
## We have a classification task

- How many targets/classes are there?
- What is the current input (features)?
- What is the current activation function?
- Is the problem nonlinear? Do we have nonlinearities in our network?
- Will the network correctly separate the classes with the current parameters?
- Run the example as given





# Let's put everything we have learned to practice!

Go to: <https://playground.tensorflow.org/>

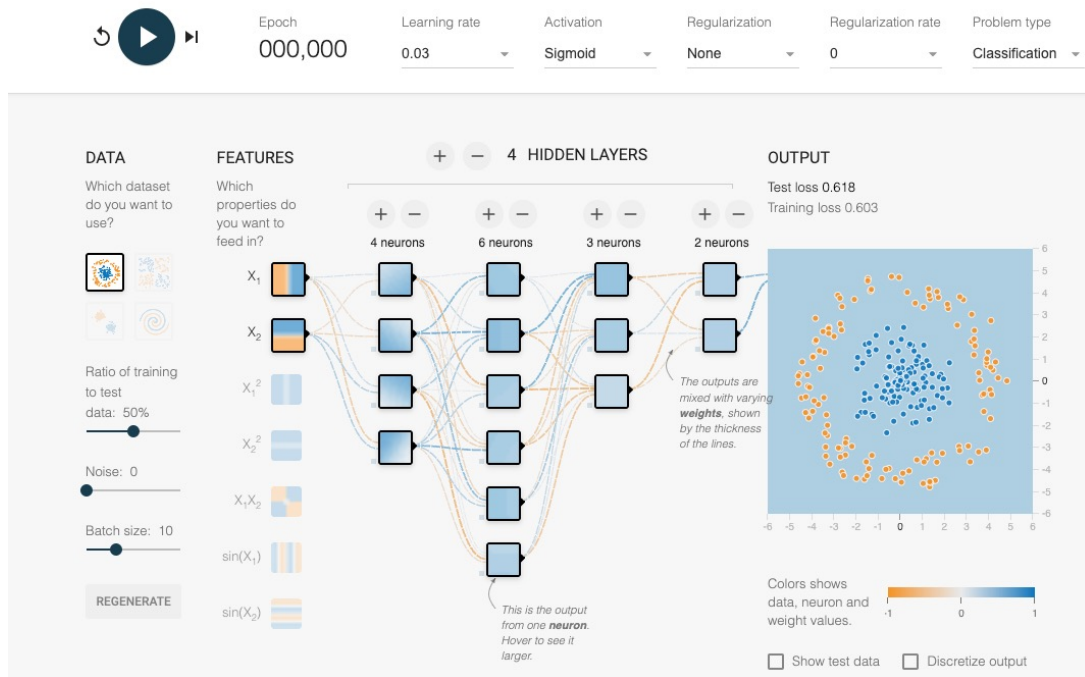


- Change the activation function to linear and run the example. What happens?
- Try adding some nonlinear combination of features. What happens? Which one works?
- Go back to only linear features. Compare the convergence speed of Tanh and sigmoid activation functions. Why is one faster than the other? Try then ReLU.



# Let's put everything we have learned to practice!

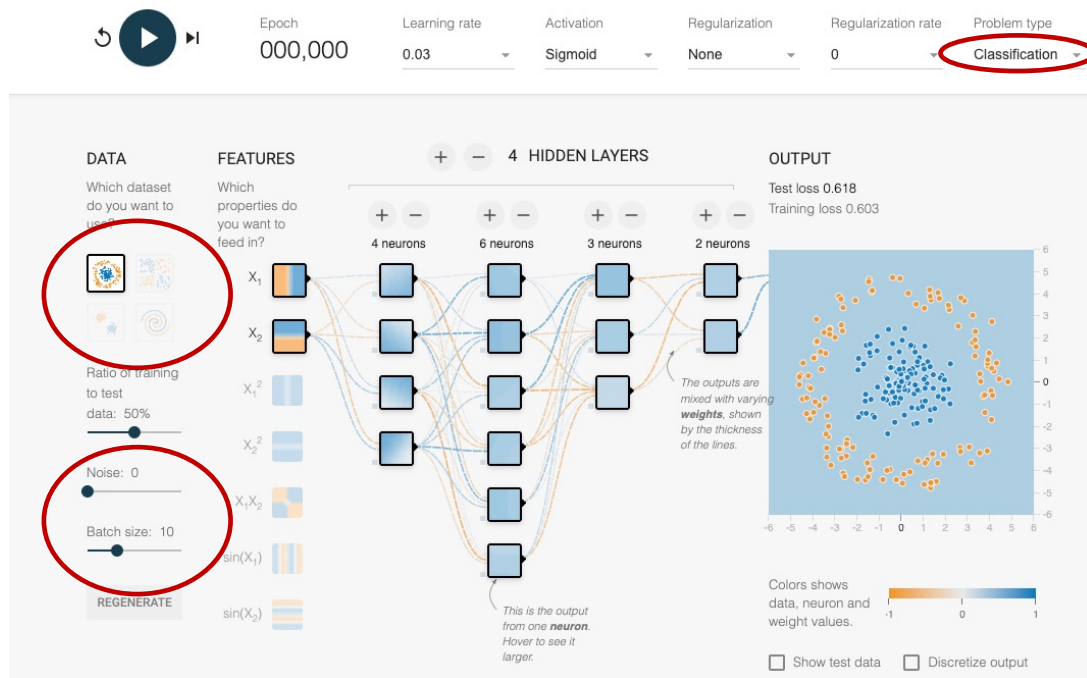
Go to: <https://playground.tensorflow.org/>



- Increase the size of the network by adding a pair of layers.
- What happens when you try to run it with the sigmoid activation function?
- Increase the number of layers to 5, with 5 or 6 neurons per layer. Run it with the Tanh activation function. What happens to the loss?
- How can we fix it?

# Let's put everything we have learned to practice!

Go to: <https://playground.tensorflow.org/>



## Explore by yourself! (~15 min)

- Try the different classification datasets. Which features fit each problem best?
- Play with the batch size and noise.
- Try the second dataset of the regression task.

# Tutorials

<https://github.com/machine-learning-tutorial/neural-networks>

# Thank you for your attention!

## What questions do you have?

Resources:

<https://ml-cheatsheet.readthedocs.io/>

<https://notesonai.com/>

<https://buildmedia.readthedocs.org/media/pdf/ml-cheatsheet/latest/ml-cheatsheet.pdf>

<http://introtodeeplearning.com/>

<https://www.offconvex.org/>